

# Lecture 2: Instruction Semantics

---

Fall 2023  
Jason Tang

# Topics

---

- CPU operations
- Instruction types
- Functional calls

# Instructions

---

- Although there are many different CPU designs, they all share similar **instruction sets**:
  - Fundamental operations that all computers must provide
  - Designers have same goal of finding a machine language that maximizes performance while minimizing cost
- This class is based around the ARMv8-A design
  - Designed by Arm holdings, a British semiconductor company
  - Used in nearly all modern smartphones, tablets, and smartwatches

# Instruction Set Architecture (ISA)

---

- ISA: abstract interface between hardware and lowest level of software, encompassing all necessary information to write and run a machine language program
  - Includes instructions, registers, memory access, I/O, etc.
  - Enables different *implementations* of ISA to run the same software
- **Application Binary Interface** (ABI): user portion of instruction set that defines how software should use the ISA to pass data between themselves and the underlying operating system
  - Defines a standard for binary portability across computers, OS revisions, and compiler revisions

# Changes in CPU Popularity

Top 10 semiconductor companies by revenue

	1990	2000	2010	2020	2030
1	NEC	Intel	Intel	Intel	Leadership to be determined
2	Toshiba	Toshiba	Samsung	Samsung	
3	Hitachi	NEC	Toshiba	SK Hynix	
4	Intel	☆ Samsung	Texas Instruments	Micron	
5	Motorola	Texas Instruments	Renesas <sup>1</sup>	Qualcomm	
6	Fujitsu	Motorola	☆ SK Hynix	☆ Broadcom	
7	Mitsubishi	☆ STMicroelectronics	STMicroelectronics	☆ Nvidia	
8	Texas Instruments	Hitachi	☆ Micron	Texas Instruments	
9	Philips	☆ Infineon	☆ Qualcomm	☆ Apple	
10	Matsushita	Philips	Elpida <sup>2</sup>	Infineon	
	Dropped out of top 10:	<ul style="list-style-type: none"> <li>Fujitsu</li> <li>Mitsubishi</li> <li>Matsushita</li> </ul>	<ul style="list-style-type: none"> <li>Motorola</li> <li>Hitachi</li> <li>Infineon</li> <li>Philips</li> </ul>	<ul style="list-style-type: none"> <li>Renesas</li> <li>STMicroelectronics</li> <li>Elpida</li> </ul>	

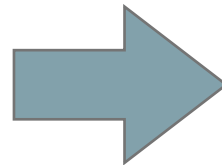
- In the past ten years, the shift has been from high performance computers to handheld and embedded systems

# Operations of Computer Hardware

---

- Assembly language is a symbolic representation of what the processor actually understands
- Different assembly language syntaxes can exist for the same ISA
  - See AT&T vs Intel debate
- In ARMv8-A, each assembly line translates to a single instruction
- Example C code and resulting assembly:

```
f = (g + h) - (i + j)
```



```
add    w1, w1, w0
add    w0, w2, w3
sub    w0, w1, w0
```

# Operands of Computer Hardware

---

- Registers are the bricks of computer construction
  - Hardware design primitives visible to programmers
- Size and number of registers differ by architectures
  - For ARMv8-A, the 31 general purpose registers are 64 bits wide
  - When referring to the lower 32 bits, use notation **wn**; to refer to all 64 bits use **xn**
- For many operations, the destination register is listed first

<b>add</b>	<b>w1</b>	<b>,</b>	<b>w1</b>	<b>,</b>	<b>w0</b>
<b>add</b>	<b>w0</b>	<b>,</b>	<b>w2</b>	<b>,</b>	<b>w3</b>
<b>sub</b>	<b>w0</b>	<b>,</b>	<b>w1</b>	<b>,</b>	<b>w0</b>

# Common Addressing Modes

---

- **Direct** (or **immediate**): value encoded in instruction
- **Register**: value stored within register
- **Register indirect**: value stored in memory at the address given by a register
- **Register displacement**: base (a register) + offset (an immediate value)
  - **Scaled displacement**:  $\text{base} + \text{offset} \times \text{scale}$
- **PC-relative**:  $\text{Program Counter} + \text{offset} \times \text{scale}$



# Complex Data Representation

---

- Given limited number of registers, remaining data (e.g., structs and arrays) are kept in memory
- Data transfer instructions to move data in memory to / from a register are traditionally called **load** / **store**
- Accessing a specific memory address is done indirectly, via register displacement
- Example in C: **g = h + A[7]**
  - Let **A** be an array of **ints** (32-bits each), and its starting address is in register w0

- Then resulting assembly is:

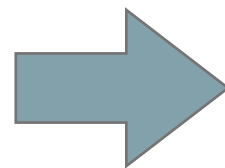
<b>ldr</b>	<b>w0, [x0, #28]</b>
<b>add</b>	<b>x2, x0, x1</b>

# Compilation and Address Formation

---

- Compiler keeps frequently used data values in registers
  - In many architectures (including ARMv8-A), arithmetic operands *must* be a register, not a location in memory
  - Thus compiler keeps frequently accessed values in registers
  - Stores everything else (including larger data structures and arrays) in main memory
- Address in memory = starting address + (index × element size)

```
int A[] = { ... };  
... A[7];
```



```
ldr    w0, [w0, #28]
```

# Byte Addresses

---

- When operating on data that is not **word aligned** within memory (e.g., an individual byte), need to specify which byte within the word to use
  - **Big-endian** (also known as network order): Motorola 68k, SuperH
  - **Little-endian**: x86-64
  - **Bi-endian**: can switch endianness as necessary
    - PowerPC starts in big-endian, can switch to little-endian
    - ARMv7 and ARMv8-A start in little-endian, can switch to big-endian

# Endian Example

---

- Example: to store the hexadecimal value 0x12345678 in memory starting at address 0x1000:

	memory cells			
	0x1000	0x1001	0x1002	0x1003
big-endian	0x12	0x34	0x56	0x78
little-endian	0x78	0x56	0x34	0x12

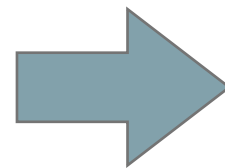
- Watch out for “**mixed-endian**” machines that may have other encoding schemes
- In ARMv8-A, **ldr** loads an entire word, while **ldrb** loads a single byte

# Stored Program Concept

---

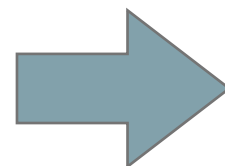
- Modern computers are built on two key principles:
  - Instructions and data are represented as numbers (specifically, in binary)
  - Instructions and data are stored in memory as numbers
- Compiler transforms code into instructions ([code generation](#)), and then the assembler translates instructions ([assembles](#)) into underlying numeric representation for that instruction set architecture

**f = (g + h) - (i + j)**



```
add    w1, w1, w0
add    w0, w2, w3
sub    w0, w1, w0
```

```
add    w1, w1, w0
add    w0, w2, w3
sub    w0, w1, w0
```



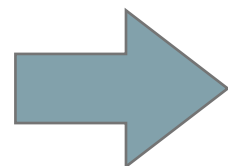
```
0x0b000021
0x0b030040
0x0b000020
```

# ARMv8-A Instruction Encoding

- In ARMv8-A, [most] instructions are encoded with exactly 32 bits ([a fixed instruction length](#))
  - x86-64 has [variable instruction length](#)
- [Instruction encoding](#) defines how instructions are numerically represented
  - Example: so-called “register-type” instructions are encoded as:

R-Type	opcode	$R_m$	shamt	$R_n$	$R_d$
	11 bits	5 bits	6 bits	5 bits	5 bits

**add w1, w1, w0**  
(0x0b000021)



00001011000	00000	000000	00001	00001
add	second operand = R0	shift amount = 0	first operand = R1	dest register = R1

# ARMv8-A Instruction Encoding

---

- Data transfers (loads and stores) use similar encoding
  - Offset must be from  $-256$  to  $+255$  (a range of  $2^9$  bits) of base register

D-Type	opcode	address	op2	$R_n$	$R_t$
	11 bits	9 bits	2 bits	5 bits	5 bits

- Immediate instructions contain a direct value instead of a register number
  - Immediate value limited to 12 bits

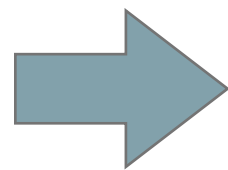
I-Type	opcode	immediate	$R_n$	$R_d$
	10 bits	12 bits	5 bits	5 bits

# Branching

---

- Several ways to change program flow, like C's **if** and **goto** statements
- Use **b** to unconditionally jump to a **label**
- Use **cbz** (compare and branch if zero) and **cbnz** (compare and branch if not zero) to conditionally jump to a label

```
if (x == 0)
    x = 10
else
    x += 20
```



```
cbnz    x0, else
mov     x0, #10
b       end_if
else:
    add     x0, x0, #20
end_if:
    /* end of if stmt */
```



# Condition Codes

---

- Many instructions can affect **condition codes**:

condition bit	set when most recent instruction resulted in:
zero (Z)	equal
carry (C)	carry out of most significant bit
negative (N)	most significant bit was 1
overflow (V)	an overflow

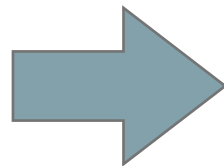
- Use a combination of these codes to express full range of comparisons: less than (**lt**), greater than (**gt**), less than or equal (**le**), greater than or equal (**ge**), equal (**eq**), and not equal (**ne**)
  - Example: **b.lt** to branch if most recent condition resulted in less than

# Setting Condition Codes

---

- Many ARMv8-A arithmetic instructions can have a **s** suffix to set condition code
- Can also directly compare a register to an immediate value via **cmp** instruction
- Example: calculate  $3^x$  using a loop

```
result = 1;
while (x > 0) {
    result *= 3;
    x--;
}
```



```
start:
    mov     x0, #1
    mov     x2, #3
    cbz     x1, end
    mul     x0, x0, x2
    subs    x1, x1, #1
    b.gt    start
end:
    /* end of while */
```

# Function Calls

---

- Calling a function depends upon the **calling convention** for the language, compiler, and instruction set
  - Need to agree on where store function arguments (register/stack/memory), where to store resulting value, and which registers are to be preserved by the **caller** and which by the **callee**
- When updating the calling stack, need to decide who updates the stack pointer (caller or callee) and by how much (**prologue** and **epilogue**)
  - For many architectures, the stack must be aligned to some boundary (every 4 bytes or every 8 bytes)

# ARMv8-A Procedure Call Standard

---

- First eight arguments to function are stored in registers **x0** through **x7**
  - Use stack if more space is needed for incoming parameters
- **x0** holds the result of the function
- **x19** through **x29** are callee-saved
- **x30** is the **link register**
- **SP** is the stack register (special register 31)
- All other registers are caller-saved

# Function Prologue and Epilogue

- Callee is responsible to moving stack pointer to hold local variables (prologue) and then cleaning up after itself (epilogue)
- Prologue and epilogue are added by the compiler; they are not hardware enforced

```
static int __attribute__((noinline))  
x(int g, int h, int i, int j) {  
    int f = (g + h) - (i + j);  
    return f;  
}
```



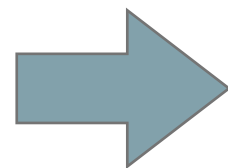
<b>x:</b>			Prologue
sub	sp,	sp, #0x20	
str	w0,	[sp, #12]	
str	w1,	[sp, #8]	
str	w2,	[sp, #4]	
str	w3,	[sp]	
ldr	w1,	[sp, #12]	
ldr	w0,	[sp, #8]	
add	w1,	w1, w0	
ldr	w2,	[sp, #4]	
ldr	w0,	[sp]	
add	w0,	w2, w0	Epilogue
sub	w0,	w1, w0	
str	w0,	[sp, #28]	
ldr	w0,	[sp, #28]	
add	sp,	sp, #0x20	
ret			

# Calling Functions

---

- Architectures vary on how to call and return from a function, specifically on where to store/load the return address
  - On x86-64, return address is pushed onto the stack upon **call**; **ret** instruction pops top-most value from stack and jumps to that address
  - On ARMv8-A, return address is stored in **x30** as a side-effect of **bl**; **ret** instruction jumps to the address stored in **x30**
- Callee is responsible for saving return address if it is not a **leaf function**

```
some_func(10, 20, 30, 40);
```



```
mov    x0, #0x0a
mov    x1, #0x14
mov    x2, #0x1e
mov    x3, #0x28
bl     some_func
```